

ASSIGNMENT 4

21BCE8842

ANIKET SAXENA

1.Download the Employee Attrition Dataset

<https://www.kaggle.com/datasets/patelprashant/employee-attrition>

2.Perform Data Preprocessing 3.Model Building using Logistic Regression and Decision Tree 4.Calculate Performance metrics

Data Collection.

- o Collect the dataset or Create the dataset

Data Preprocessing.

- o Import the Libraries.
- o Importing the dataset.
- o Checking for Null Values.
- o Data Visualization.
- o Outlier Detection
- o Splitting Dependent and Independent variables
- o Encoding
- o Feature Scaling.
- o Splitting Data into Train and Test.

Model Building

- o Import the model building Libraries
- o Initializing the model
- o Training and testing the model
- o Evaluation of Model
- o Save the Model

Data Preprocessing

- o Import the Libraries
- o Importing the Dataset
- o Handling Null values
- o Data Visualization
- o Outlier Detection
- o Splitting Dependent and Independent variable
- o Encoding
- o Feature Scaling
- o Splitting Data in Train and Test

Import the Libraries

```
In [1]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns
```

Importing the Dataset

```
In [2]: df=pd.read_csv('WA_Fn-UseC_-HR-Employee-Attrition.csv')
```

```
In [3]: df.head()
```

```
Out[3]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education
0	41	Yes	Travel_Rarely	1102	Sales	1	2
1	49	No	Travel_Frequently	279	Research & Development	8	1
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2
3	33	No	Travel_Frequently	1392	Research & Development	3	4
4	27	No	Travel_Rarely	591	Research & Development	2	1

5 rows × 35 columns

```
In [4]: df.shape
```

```
Out[4]: (1470, 35)
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Age              1470 non-null    int64  
 1   Attrition        1470 non-null    object  
 2   BusinessTravel   1470 non-null    object  
 3   DailyRate         1470 non-null    int64  
 4   Department        1470 non-null    object  
 5   DistanceFromHome 1470 non-null    int64  
 6   Education         1470 non-null    int64  
 7   EducationField    1470 non-null    object  
 8   EmployeeCount     1470 non-null    int64  
 9   EmployeeNumber    1470 non-null    int64  
 10  EnvironmentSatisfaction 1470 non-null    int64  
 11  Gender            1470 non-null    object  
 12  HourlyRate        1470 non-null    int64  
 13  JobInvolvement   1470 non-null    int64  
 14  JobLevel          1470 non-null    int64  
 15  JobRole           1470 non-null    object  
 16  JobSatisfaction   1470 non-null    int64  
 17  MaritalStatus     1470 non-null    object  
 18  MonthlyIncome     1470 non-null    int64  
 19  MonthlyRate       1470 non-null    int64  
 20  NumCompaniesWorked 1470 non-null    int64  
 21  Over18            1470 non-null    object  
 22  OverTime          1470 non-null    object  
 23  PercentSalaryHike 1470 non-null    int64  
 24  PerformanceRating 1470 non-null    int64  
 25  RelationshipSatisfaction 1470 non-null    int64  
 26  StandardHours     1470 non-null    int64  
 27  StockOptionLevel   1470 non-null    int64  
 28  TotalWorkingYears 1470 non-null    int64  
 29  TrainingTimesLastYear 1470 non-null    int64  
 30  WorkLifeBalance   1470 non-null    int64  
 31  YearsAtCompany    1470 non-null    int64  
 32  YearsInCurrentRole 1470 non-null    int64  
 33  YearsSinceLastPromotion 1470 non-null    int64  
 34  YearsWithCurrManager 1470 non-null    int64  
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
```

In [6]: `df.describe()`

Out[6]:

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeCount	EmployeeCount
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.0	1470.0
mean	36.923810	802.485714	9.192517	2.912925	1.0	1.0	1.0
std	9.135373	403.509100	8.106864	1.024165	0.0	0.0	0.0
min	18.000000	102.000000	1.000000	1.000000	1.0	1.0	1.0
25%	30.000000	465.000000	2.000000	2.000000	1.0	1.0	1.0
50%	36.000000	802.000000	7.000000	3.000000	1.0	1.0	1.0
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1.0	1.0
max	60.000000	1499.000000	29.000000	5.000000	1.0	2.0	2.0

8 rows × 26 columns



Handling Null values

In [7]: `df.isnull().any()`

```
Out[7]: Age           False
Attrition      False
BusinessTravel  False
DailyRate       False
Department     False
DistanceFromHome False
Education       False
EducationField  False
EmployeeCount   False
EmployeeNumber  False
EnvironmentSatisfaction False
Gender          False
HourlyRate      False
JobInvolvement  False
JobLevel         False
JobRole          False
JobSatisfaction False
MaritalStatus   False
MonthlyIncome    False
MonthlyRate     False
NumCompaniesWorked False
Over18          False
OverTime         False
PercentSalaryHike False
PerformanceRating False
RelationshipSatisfaction False
StandardHours   False
StockOptionLevel False
TotalWorkingYears False
TrainingTimesLastYear False
WorkLifeBalance  False
YearsAtCompany   False
YearsInCurrentRole False
YearsSinceLastPromotion False
YearsWithCurrManager False
dtype: bool
```

```
In [8]: df.isnull().sum() # no null values found
```

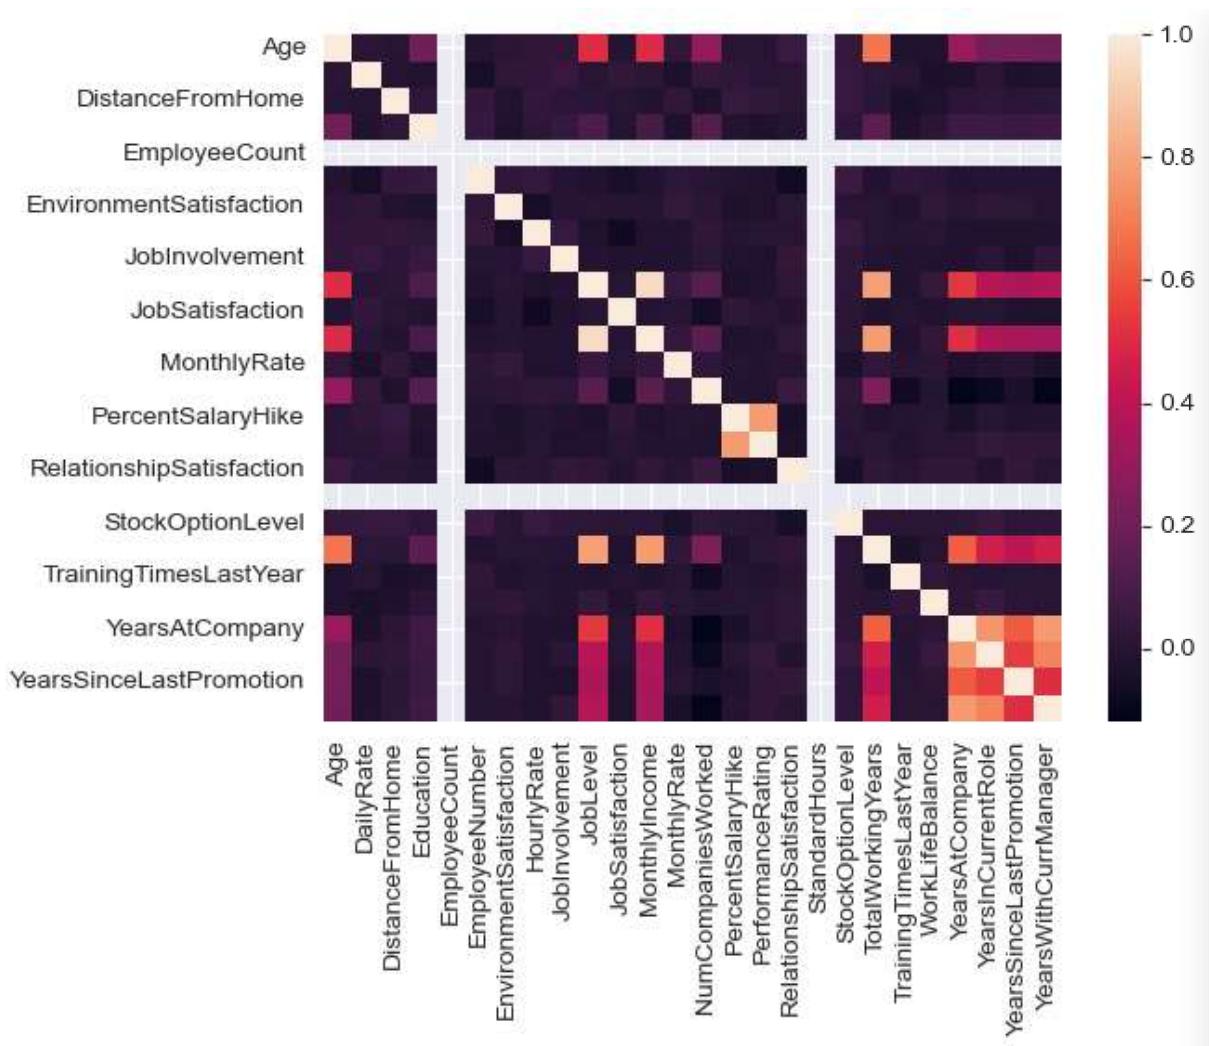
```
Out[8]: Age          0  
Attrition      0  
BusinessTravel  0  
DailyRate       0  
Department     0  
DistanceFromHome 0  
Education       0  
EducationField   0  
EmployeeCount    0  
EmployeeNumber   0  
EnvironmentSatisfaction 0  
Gender          0  
HourlyRate      0  
JobInvolvement   0  
JobLevel         0  
JobRole          0  
JobSatisfaction  0  
MaritalStatus    0  
MonthlyIncome     0  
MonthlyRate       0  
NumCompaniesWorked 0  
Over18           0  
OverTime          0  
PercentSalaryHike 0  
PerformanceRating 0  
RelationshipSatisfaction 0  
StandardHours     0  
StockOptionLevel   0  
TotalWorkingYears  0  
TrainingTimesLastYear 0  
WorkLifeBalance    0  
YearsAtCompany     0  
YearsInCurrentRole 0  
YearsSinceLastPromotion 0  
YearsWithCurrManager 0  
dtype: int64
```

Data Visualization

```
In [9]: sns.heatmap(df.corr())
```

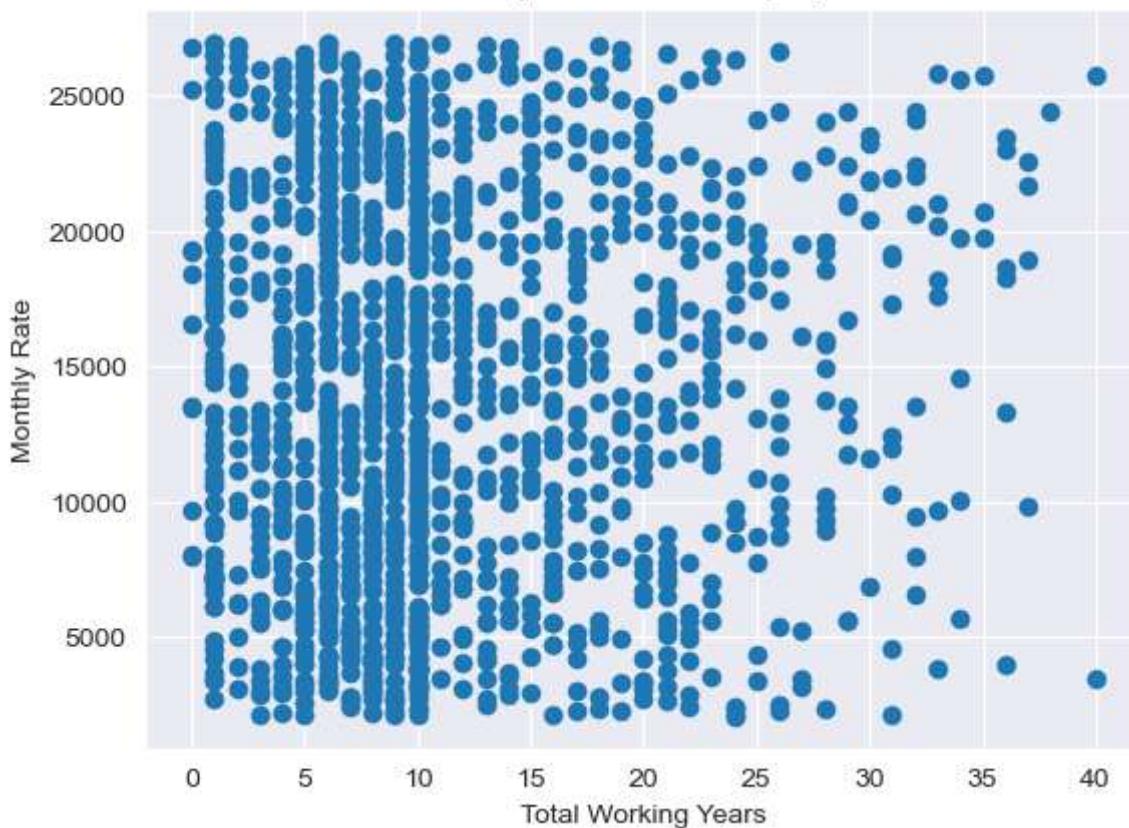
C:\Users\DELL\AppData\Local\Temp\ipykernel_2984\1072140413.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.
sns.heatmap(df.corr())

```
Out[9]: <Axes: >
```



```
In [10]: plt.scatter(df['TotalWorkingYears'], df['MonthlyRate'])
plt.title("Comparison of Total Working Years against \n Monthly Rate for all Employ
plt.xlabel("Total Working Years")
plt.ylabel("Monthly Rate")
plt.show()
```

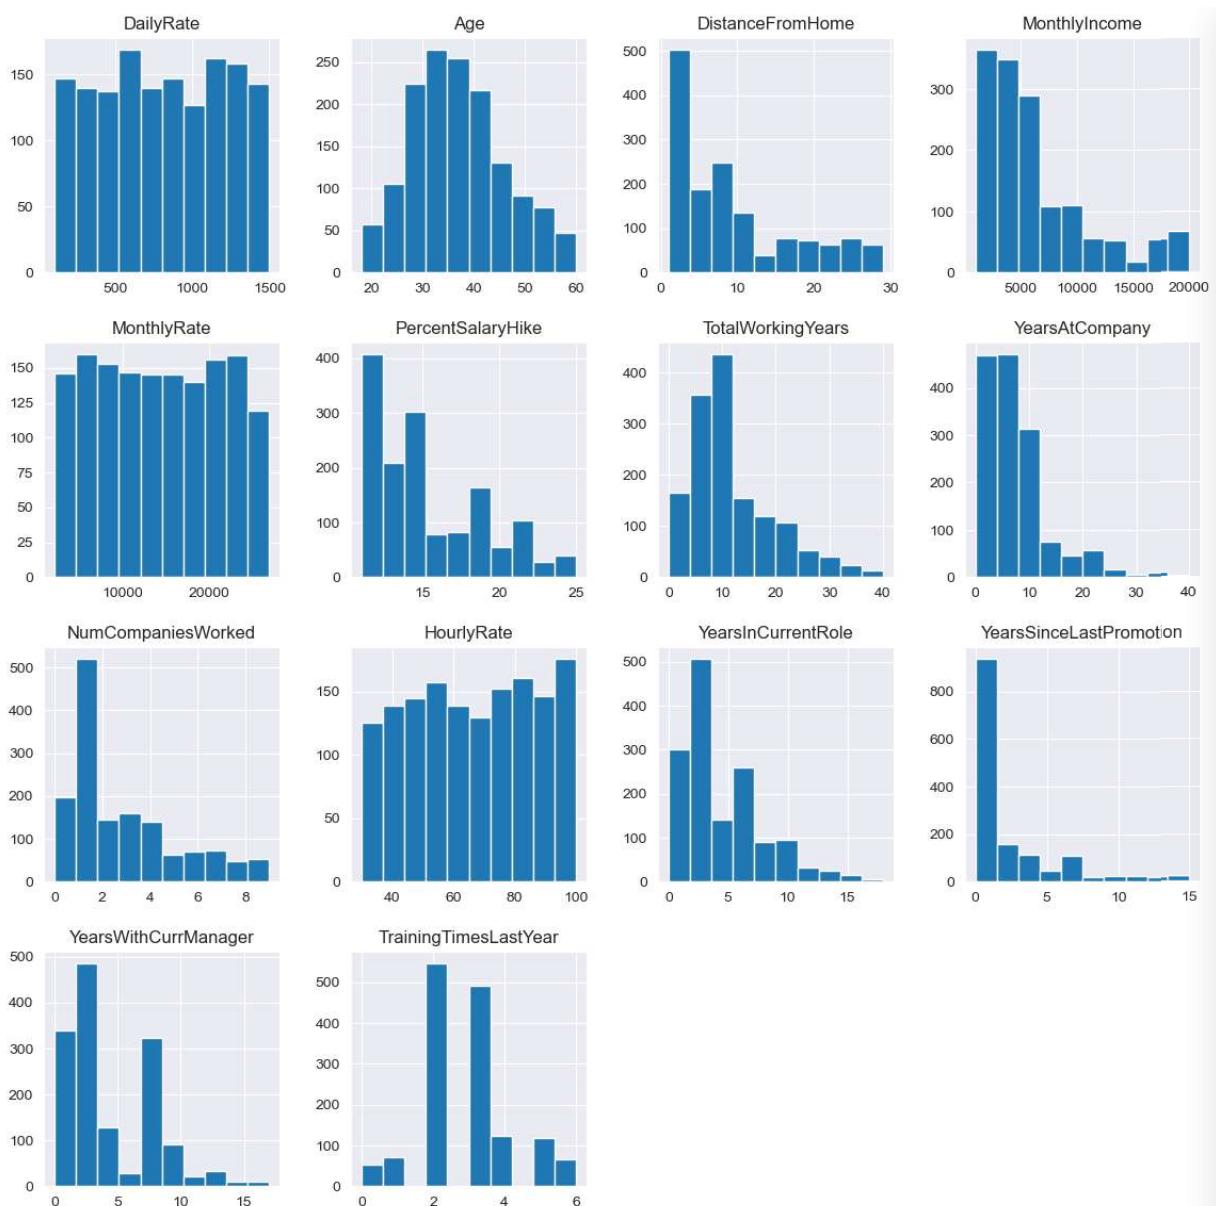
Comparison of Total Working Years against Monthly Rate for all Employees



```
In [11]: numeric_columns = ['DailyRate', 'Age', 'DistanceFromHome', 'MonthlyIncome', 'MonthlyRate',  
'YearsAtCompany', 'NumCompaniesWorked', 'HourlyRate',  
'YearsInCurrentRole', 'YearsSinceLastPromotion', 'YearsWithCurrMan
```

```
In [12]: df[numeric_columns].hist(figsize=(14, 14))
```

```
Out[12]: array([[<Axes: title={'center': 'DailyRate'}>,  
    <Axes: title={'center': 'Age'}>,  
    <Axes: title={'center': 'DistanceFromHome'}>,  
    <Axes: title={'center': 'MonthlyIncome'}>],  
   [<Axes: title={'center': 'MonthlyRate'}>,  
    <Axes: title={'center': 'PercentSalaryHike'}>,  
    <Axes: title={'center': 'TotalWorkingYears'}>,  
    <Axes: title={'center': 'YearsAtCompany'}>],  
   [<Axes: title={'center': 'NumCompaniesWorked'}>,  
    <Axes: title={'center': 'HourlyRate'}>,  
    <Axes: title={'center': 'YearsInCurrentRole'}>,  
    <Axes: title={'center': 'YearsSinceLastPromotion'}>],  
   [<Axes: title={'center': 'YearsWithCurrManager'}>,  
    <Axes: title={'center': 'TrainingTimesLastYear'}>, <Axes: >,  
    <Axes: >]], dtype=object)
```



Outlier detection

```
In [13]: df.corr()
```

```
C:\Users\DELL\AppData\Local\Temp\ipykernel_2984\592014893.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.
df.corr()
```

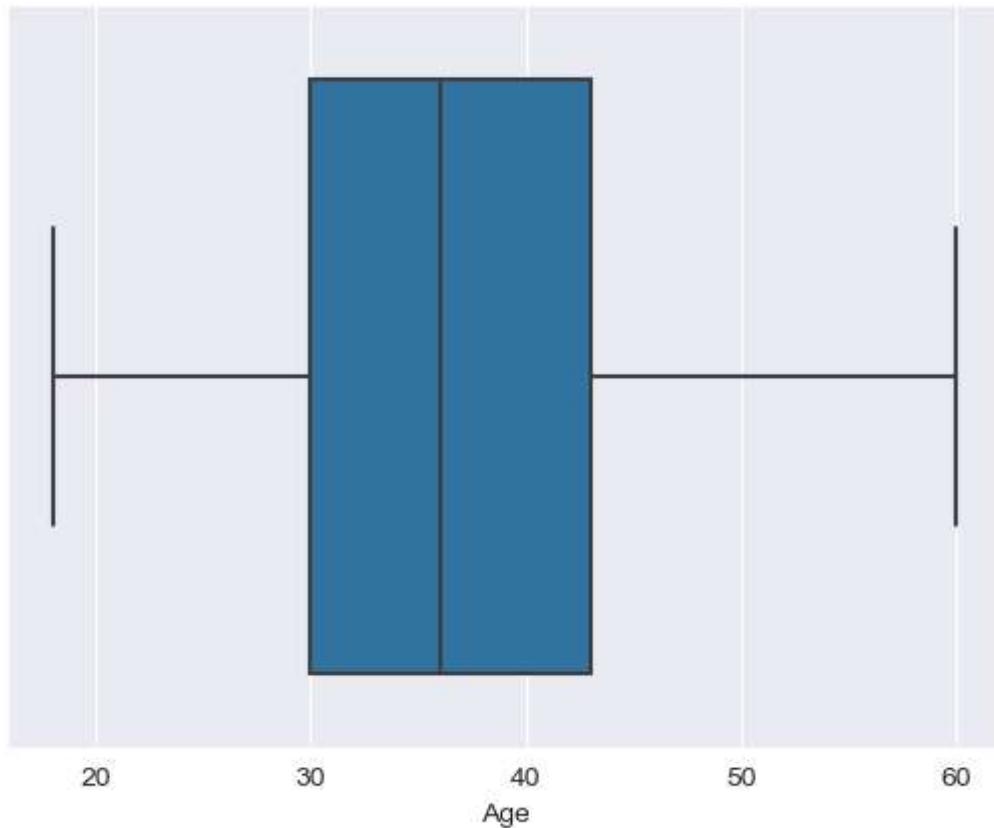
Out[13]:

	Age	DailyRate	DistanceFromHome	Education	EmployeeCo
Age	1.000000	0.010661	-0.001686	0.208034	NaN
DailyRate	0.010661	1.000000	-0.004985	-0.016806	NaN
DistanceFromHome	-0.001686	-0.004985	1.000000	0.021042	NaN
Education	0.208034	-0.016806	0.021042	1.000000	NaN
EmployeeCount	NaN	NaN	NaN	NaN	NaN
EmployeeNumber	-0.010145	-0.050990	0.032916	0.042070	NaN
EnvironmentSatisfaction	0.010146	0.018355	-0.016075	-0.027128	NaN
HourlyRate	0.024287	0.023381	0.031131	0.016775	NaN
JobInvolvement	0.029820	0.046135	0.008783	0.042438	NaN
JobLevel	0.509604	0.002966	0.005303	0.101589	NaN
JobSatisfaction	-0.004892	0.030571	-0.003669	-0.011296	NaN
MonthlyIncome	0.497855	0.007707	-0.017014	0.094961	NaN
MonthlyRate	0.028051	-0.032182	0.027473	-0.026084	NaN
NumCompaniesWorked	0.299635	0.038153	-0.029251	0.126317	NaN
PercentSalaryHike	0.003634	0.022704	0.040235	-0.011111	NaN
PerformanceRating	0.001904	0.000473	0.027110	-0.024539	NaN
RelationshipSatisfaction	0.053535	0.007846	0.006557	-0.009118	NaN
StandardHours	NaN	NaN	NaN	NaN	NaN
StockOptionLevel	0.037510	0.042143	0.044872	0.018422	NaN
TotalWorkingYears	0.680381	0.014515	0.004628	0.148280	NaN
TrainingTimesLastYear	-0.019621	0.002453	-0.036942	-0.025100	NaN
WorkLifeBalance	-0.021490	-0.037848	-0.026556	0.009819	NaN
YearsAtCompany	0.311309	-0.034055	0.009508	0.069114	NaN
YearsInCurrentRole	0.212901	0.009932	0.018845	0.060236	NaN
YearsSinceLastPromotion	0.216513	-0.033229	0.010029	0.054254	NaN
YearsWithCurrManager	0.202089	-0.026363	0.014406	0.069065	NaN

26 rows × 26 columns

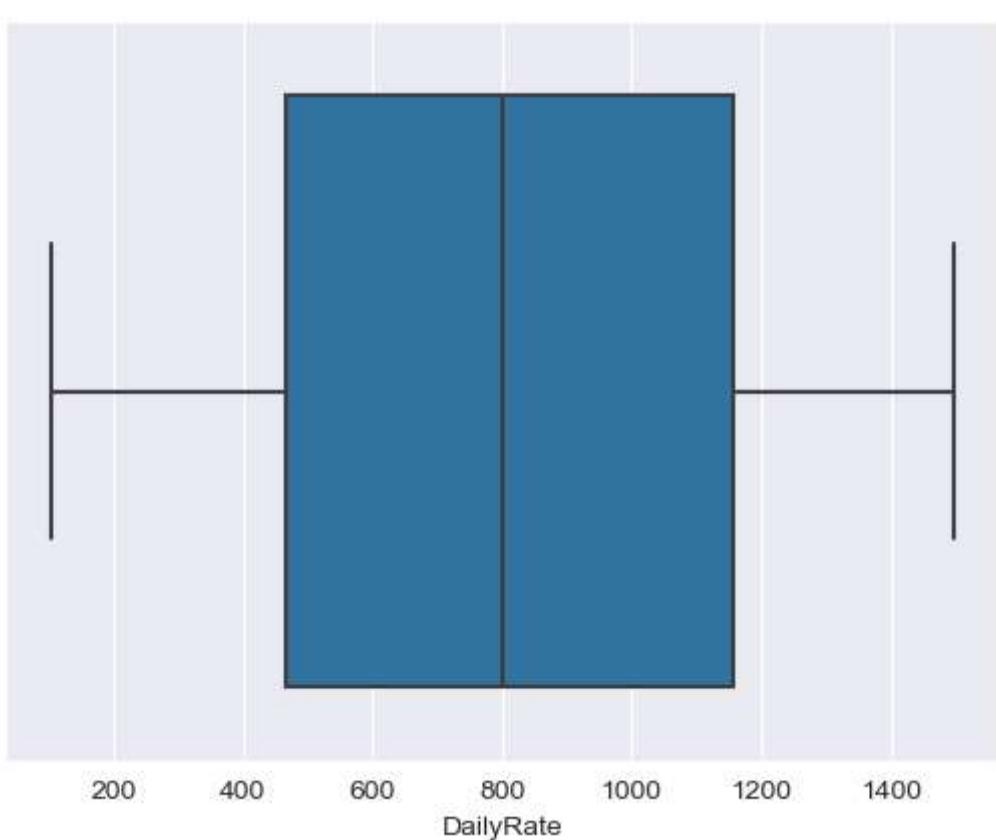
In [14]: `sns.boxplot(x='Age', data=df)`

```
Out[14]: <Axes: xlabel='Age'>
```



```
In [15]: sns.boxplot(x='DailyRate',data=df)
```

```
Out[15]: <Axes: xlabel='DailyRate'>
```



Label Encoding

```
In [16]: from sklearn.preprocessing import LabelEncoder  
le=LabelEncoder()
```

```
In [17]: columns = ['Attrition','OverTime','BusinessTravel', 'Department','Education',  
                 'EducationField','JobSatisfaction','EnvironmentSatisfaction'  
                 'StockOptionLevel','Gender', 'PerformanceRating', 'JobInvolv  
                 'JobLevel', 'JobRole', 'MaritalStatus','RelationshipSatisfac
```

```
In [18]: df1=df.copy()
```

```
In [19]: df1[columns]=df1[columns].apply(le.fit_transform)
```

```
In [20]: df1.head()
```

Out[20]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education
0	41	1	2	1102	2	1	1
1	49	0	1	279	1	8	0
2	37	1	2	1373	1	2	1
3	33	0	1	1392	1	3	3
4	27	0	2	591	1	2	0

5 rows × 35 columns

In [21]: df1.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Age              1470 non-null    int64  
 1   Attrition        1470 non-null    int32  
 2   BusinessTravel   1470 non-null    int32  
 3   DailyRate         1470 non-null    int64  
 4   Department        1470 non-null    int32  
 5   DistanceFromHome 1470 non-null    int64  
 6   Education         1470 non-null    int64  
 7   EducationField    1470 non-null    int32  
 8   EmployeeCount     1470 non-null    int64  
 9   EmployeeNumber    1470 non-null    int64  
 10  EnvironmentSatisfaction 1470 non-null    int64  
 11  Gender            1470 non-null    int32  
 12  HourlyRate        1470 non-null    int64  
 13  JobInvolvement   1470 non-null    int64  
 14  JobLevel          1470 non-null    int64  
 15  JobRole           1470 non-null    int32  
 16  JobSatisfaction   1470 non-null    int64  
 17  MaritalStatus     1470 non-null    int32  
 18  MonthlyIncome     1470 non-null    int64  
 19  MonthlyRate       1470 non-null    int64  
 20  NumCompaniesWorked 1470 non-null    int64  
 21  Over18            1470 non-null    int32  
 22  OverTime          1470 non-null    int32  
 23  PercentSalaryHike 1470 non-null    int64  
 24  PerformanceRating 1470 non-null    int64  
 25  RelationshipSatisfaction 1470 non-null    int64  
 26  StandardHours     1470 non-null    int64  
 27  StockOptionLevel   1470 non-null    int64  
 28  TotalWorkingYears  1470 non-null    int64  
 29  TrainingTimesLastYear 1470 non-null    int64  
 30  WorkLifeBalance   1470 non-null    int64  
 31  YearsAtCompany    1470 non-null    int64  
 32  YearsInCurrentRole 1470 non-null    int64  
 33  YearsSinceLastPromotion 1470 non-null    int64  
 34  YearsWithCurrManager 1470 non-null    int64  
dtypes: int32(9), int64(26)
memory usage: 350.4 KB
```

Spliting independent and dependent variable

In [22]: df1.corr()

Out[22]:

	Age	Attrition	BusinessTravel	DailyRate	Department	Dist
Age	1.000000	-0.159205	0.024751	0.010661	-0.031882	
Attrition	-0.159205	1.000000	0.000074	-0.056652	0.063991	
BusinessTravel	0.024751	0.000074	1.000000	-0.004086	-0.009044	
DailyRate	0.010661	-0.056652	-0.004086	1.000000	0.007109	
Department	-0.031882	0.063991	-0.009044	0.007109	1.000000	
DistanceFromHome	-0.001686	0.077924	-0.024469	-0.004985	0.017225	
Education	0.208034	-0.031373	0.000757	-0.016806	0.007996	
EducationField	-0.040873	0.026846	0.023724	0.037709	0.013720	
EmployeeCount	NaN	NaN	NaN	NaN	NaN	
EmployeeNumber	-0.010145	-0.010577	-0.015578	-0.050990	-0.010895	
EnvironmentSatisfaction	0.010146	-0.103369	0.004174	0.018355	-0.019395	
Gender	-0.036311	0.029453	-0.032981	-0.011716	-0.041583	
HourlyRate	0.024287	-0.006846	0.026528	0.023381	-0.004144	
JobInvolvement	0.029820	-0.130016	0.039062	0.046135	-0.024586	
JobLevel	0.509604	-0.169105	0.019311	0.002966	0.101963	
JobRole	-0.122427	0.067151	0.002724	-0.009472	0.662431	
JobSatisfaction	-0.004892	-0.103481	-0.033962	0.030571	0.021001	
MaritalStatus	-0.095029	0.162070	0.024001	-0.069586	0.056073	
MonthlyIncome	0.497855	-0.159840	0.034319	0.007707	0.053130	
MonthlyRate	0.028051	0.015170	-0.014107	-0.032182	0.023642	
NumCompaniesWorked	0.299635	0.043494	0.020875	0.038153	-0.035882	
Over18	NaN	NaN	NaN	NaN	NaN	
OverTime	0.028062	0.246118	0.016543	0.009135	0.007481	
PercentSalaryHike	0.003634	-0.013478	-0.029377	0.022704	-0.007840	
PerformanceRating	0.001904	0.002889	-0.026341	0.000473	-0.024604	
RelationshipSatisfaction	0.053535	-0.045872	-0.035986	0.007846	-0.022414	
StandardHours	NaN	NaN	NaN	NaN	NaN	
StockOptionLevel	0.037510	-0.137145	-0.016727	0.042143	-0.012193	
TotalWorkingYears	0.680381	-0.171063	0.034226	0.014515	-0.015762	
TrainingTimesLastYear	-0.019621	-0.059478	0.015240	0.002453	0.036875	

	Age	Attrition	BusinessTravel	DailyRate	Department	Dist
WorkLifeBalance	-0.021490	-0.063939	-0.011256	-0.037848	0.026383	
YearsAtCompany	0.311309	-0.134392	-0.014575	-0.034055	0.022920	
YearsInCurrentRole	0.212901	-0.160545	-0.011497	0.009932	0.056315	
YearsSinceLastPromotion	0.216513	-0.033019	-0.032591	-0.033229	0.040061	
YearsWithCurrManager	0.202089	-0.156199	-0.022636	-0.026363	0.034282	

35 rows × 35 columns

```
In [23]: x = df1.drop(columns='Attrition')
```

```
In [24]: y = df1['Attrition']
```

```
In [25]: x.head()
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationF
0	41	2	1102	2		1	1
1	49	1	279	1		8	0
2	37	2	1373	1		2	1
3	33	1	1392	1		3	3
4	27	2	591	1		2	0

5 rows × 34 columns

```
In [26]: y.head()
```

```
Out[26]: 0    1
1    0
2    1
3    0
4    0
Name: Attrition, dtype: int32
```

Feature Scaling

```
In [27]: from sklearn.preprocessing import MinMaxScaler
```

```
In [28]: ms=MinMaxScaler()
ms
```

```
Out[28]: ▾ MinMaxScaler
```

```
    MinMaxScaler()
```

```
In [29]: x_scaled=pd.DataFrame(ms.fit_transform(x),columns=x.columns)
```

```
In [30]: x_scaled
```

```
Out[30]:
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Ed
0	0.547619		1.0	0.715820	1.0	0.000000	0.25
1	0.738095		0.5	0.126700	0.5	0.250000	0.00
2	0.452381		1.0	0.909807	0.5	0.035714	0.25
3	0.357143		0.5	0.923407	0.5	0.071429	0.75
4	0.214286		1.0	0.350036	0.5	0.035714	0.00
...
1465	0.428571		0.5	0.559771	0.5	0.785714	0.25
1466	0.500000		1.0	0.365784	0.5	0.178571	0.00
1467	0.214286		1.0	0.037938	0.5	0.107143	0.50
1468	0.738095		0.5	0.659270	1.0	0.035714	0.50
1469	0.380952		1.0	0.376521	0.5	0.250000	0.50

1470 rows × 34 columns



Splitting Data into Train and test Split

```
In [31]: from sklearn.model_selection import train_test_split
```

```
In [32]: x_train,x_test,y_train,y_test=train_test_split(x_scaled,y,test_size=0.2,random_stat
```

```
In [33]: x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
Out[33]: ((1176, 34), (294, 34), (1176,), (294,))
```

• Model Building

- o Import the model building Libraries
- o Initializing the model

- o Training and testing the model
 - o Evaluation of Model
 - o Save the Model

LOGISTIC REGRESSION

```
In [34]: from sklearn.linear_model import LogisticRegression
```

```
In [35]: model_logistic=LogisticRegression()  
model_logistic
```

Out[35]: `LogisticRegression`
`LogisticRegression()`

```
In [36]: model logistic.fit(x_train,y_train)
```

Out[36]: `LogisticRegression`
`LogisticRegression()`

```
In [37]: pred_logistic=model_logistic.predict(x_test)
```

```
In [38]: pred logistic
```

```
In [39]: y_test
```

```
Out[39]:
```

442	0
1091	0
981	1
785	0
1332	1
	..
1439	0
481	0
124	1
198	0
1229	0

```
Name: Attrition, Length: 294, dtype: int32
```

DECISION TREE CLASSIFIER

```
In [41]: from sklearn.tree import DecisionTreeClassifier
```

```
In [42]: model_decision_tree=DecisionTreeClassifier()  
model decision tree
```

```
Out[42]: ▾ DecisionTreeClassifier  
DecisionTreeClassifier()
```

```
In [43]: model decision tree.fit(x_train,y_train)
```

```
Out[43]: ▾ DecisionTreeClassifier  
DecisionTreeClassifier()
```

```
In [44]: pred = model.decision_tree.predict(x_test)
```

```
In [45]: pred decision tree
```

```
In [46]: y_test
```

```
Out[46]: 442      0
        1091     0
        981      1
        785      0
        1332     1
        ..
       1439      0
       481      0
       124      1
       198      0
       1229     0
Name: Attrition, Length: 294, dtype: int32
```

Evaluation of Classification model

For Logistic Regression Model

```
In [47]: #Accuracy score
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report,r
```



```
In [48]: accuracy_score_logistic=accuracy_score(y_test,pred_logistic)
accuracy_score_logistic
```



```
Out[48]: 0.8843537414965986
```

```
In [49]: confusion_matrix_logistic=confusion_matrix(y_test,pred_logistic)
confusion_matrix_logistic
```



```
Out[49]: array([[242,    3],
                 [ 31,   18]], dtype=int64)
```

```
In [50]: crosstab_logistic=pd.crosstab(y_test,pred_logistic)
crosstab_logistic
```



```
Out[50]:   col_0    0    1
Attrition
          0  242    3
          1   31   18
```

```
In [51]: print(classification_report(y_test,pred_logistic))
```

	precision	recall	f1-score	support
0	0.89	0.99	0.93	245
1	0.86	0.37	0.51	49
accuracy			0.88	294
macro avg	0.87	0.68	0.72	294
weighted avg	0.88	0.88	0.86	294

```
In [52]: probability_logistic=model_logistic.predict_proba(x_test)[:,1]
```

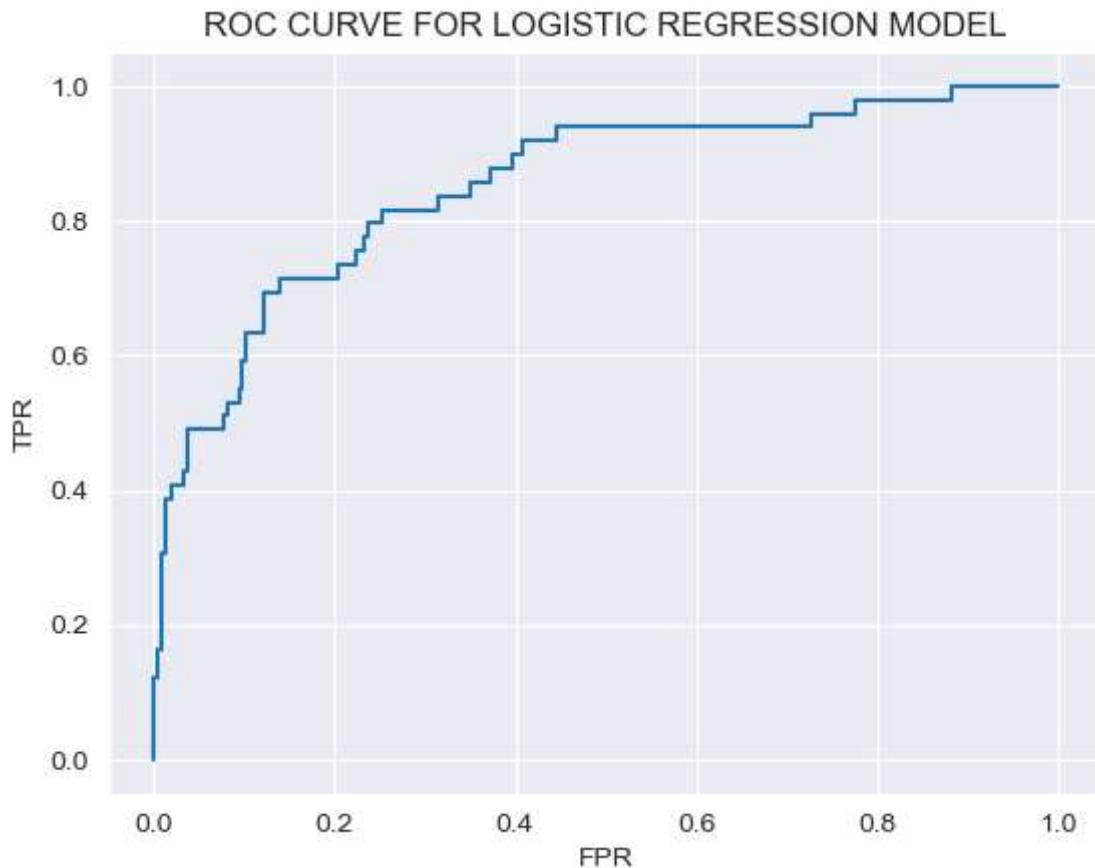
```
In [53]: probability_logistic
```

```
Out[53]: array([0.16000127, 0.20600667, 0.31532384, 0.09242886, 0.63667551,  
    0.06153061, 0.61819432, 0.0757087 , 0.00841372, 0.3912069 ,  
    0.05398439, 0.33293123, 0.02020698, 0.67215483, 0.19786547,  
    0.03454902, 0.11043981, 0.17101703, 0.04477777, 0.22783614,  
    0.2335018 , 0.01553905, 0.06464492, 0.05029956, 0.58792413,  
    0.44849464, 0.07412714, 0.04460935, 0.67666632, 0.0584383 ,  
    0.01599026, 0.03521098, 0.06963085, 0.17397462, 0.07830857,  
    0.04288032, 0.08150424, 0.07106342, 0.03622137, 0.05223965,  
    0.04862098, 0.02091497, 0.01819361, 0.01362467, 0.02873997,  
    0.50236969, 0.41553218, 0.00306874, 0.73976412, 0.51382382,  
    0.09637213, 0.48845516, 0.08036228, 0.25757243, 0.66516772,  
    0.26308027, 0.01964858, 0.30198497, 0.02919946, 0.16038964,  
    0.02102747, 0.21670232, 0.13981568, 0.0358316 , 0.37208403,  
    0.03002317, 0.29091186, 0.16041142, 0.10437497, 0.08695177,  
    0.08217589, 0.30984518, 0.08531362, 0.07420689, 0.12268651,  
    0.06192552, 0.04640904, 0.07624712, 0.19738483, 0.03236316,  
    0.00884439, 0.0244108 , 0.13635803, 0.0260104 , 0.03341008,  
    0.08186888, 0.00499397, 0.03474852, 0.03858027, 0.14602694,  
    0.26167665, 0.16667357, 0.27400109, 0.24159565, 0.02160421,  
    0.17748606, 0.34076078, 0.28022482, 0.06914126, 0.05003806,  
    0.24437761, 0.74698271, 0.35438567, 0.01920627, 0.08778845,  
    0.03255847, 0.05461351, 0.15123251, 0.06843702, 0.13752637,  
    0.09584388, 0.04669882, 0.02493091, 0.15383171, 0.07081259,  
    0.03089296, 0.0537667 , 0.11554316, 0.00881616, 0.01263271,  
    0.17552253, 0.05045234, 0.08823238, 0.82995757, 0.03017756,  
    0.0236819 , 0.0087012 , 0.1349589 , 0.16474801, 0.05202613,  
    0.01524549, 0.29278083, 0.54767448, 0.34275448, 0.04629541,  
    0.38966344, 0.61333366, 0.14552367, 0.07402366, 0.24143471,  
    0.09418418, 0.0689069 , 0.10061956, 0.19346327, 0.20026293,  
    0.03004939, 0.14900424, 0.00348846, 0.11225149, 0.15843155,  
    0.06047573, 0.18601882, 0.06085869, 0.12221317, 0.03280184,  
    0.02738799, 0.06356425, 0.08302382, 0.01541716, 0.014665 ,  
    0.38517822, 0.01264231, 0.14961974, 0.80508787, 0.11598661,  
    0.2842811 , 0.17020143, 0.1530583 , 0.02764153, 0.00613226,  
    0.04191632, 0.09782393, 0.11551417, 0.10377982, 0.01779313,  
    0.14371315, 0.10615435, 0.10298963, 0.05132621, 0.09061081,  
    0.02897383, 0.09924087, 0.00512032, 0.75108423, 0.04296968,  
    0.04062134, 0.37518972, 0.04563128, 0.7251816 , 0.10671665,  
    0.36949086, 0.38146941, 0.32095493, 0.05266802, 0.08172004,  
    0.13947833, 0.04334317, 0.01469593, 0.26413988, 0.06330966,  
    0.1614747 , 0.15380517, 0.67152357, 0.05840793, 0.27891823,  
    0.04512564, 0.46033865, 0.00348431, 0.14068967, 0.02747401,  
    0.12714133, 0.17284246, 0.07341066, 0.10099827, 0.16870885,  
    0.02560842, 0.01824031, 0.08670796, 0.02834237, 0.13710215,  
    0.08778935, 0.2200061 , 0.73401148, 0.15938978, 0.4095449 ,  
    0.01513845, 0.11306309, 0.21497506, 0.32337575, 0.03409266,  
    0.04256318, 0.32157531, 0.05454465, 0.02348479, 0.16423352,  
    0.32696147, 0.22892063, 0.00877159, 0.08198819, 0.01156361,  
    0.1408691 , 0.29235147, 0.01270305, 0.17329916, 0.04081391,  
    0.04094165, 0.42771425, 0.34958286, 0.03766772, 0.12025286,  
    0.37698923, 0.3192629 , 0.79559338, 0.05385659, 0.21597037,  
    0.06383728, 0.00570991, 0.66018187, 0.35855286, 0.37783606,  
    0.36781398, 0.03554512, 0.21718203, 0.05943622, 0.06554485,  
    0.10081475, 0.00818713, 0.26591316, 0.42809675, 0.06542835,  
    0.09296803, 0.01259826, 0.14226651, 0.05072662, 0.02372258,  
    0.02586923, 0.06760427, 0.24315648, 0.26961432, 0.19831733,
```

```
0.2652296 , 0.0165923 , 0.15784236, 0.08398982, 0.02711775,
0.18750547, 0.00783535, 0.2844239 , 0.00270742, 0.02484969,
0.22585745, 0.72775605, 0.07691547, 0.26304359])
```

```
In [54]: # Roc curve
fpr_logistic,tpr_logistic,thresholds_logistic = roc_curve(y_test,probability_logi
```

```
In [55]: plt.plot(fpr_logistic,tpr_logistic)
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('ROC CURVE FOR LOGISTIC REGRESSION MODEL')
plt.show()
```



For Decision Tree Model

```
In [56]: accuracy_score_decision=accuracy_score(y_test,pred_decision_tree)
accuracy_score_decision
```

```
Out[56]: 0.7551020408163265
```

```
In [57]: confusion_matrix_decision=confusion_matrix(y_test,pred_decision_tree)
confusion_matrix_decision
```

```
Out[57]: array([[208, 37],
 [35, 14]], dtype=int64)
```

```
In [58]: crosstab_decision=pd.crosstab(y_test,pred_decision_tree)
crosstab decision
```

Out[58]:

Attrition

```
In [59]: print(classification_report(y_test,pred_decision_tree))
```

	precision	recall	f1-score	support
0	0.86	0.85	0.85	245
1	0.27	0.29	0.28	49
accuracy			0.76	294
macro avg	0.57	0.57	0.57	294
weighted avg	0.76	0.76	0.76	294

```
In [60]: probability_decision=model_decision_tree.predict_proba(x_test)[:,1]
```

```
In [61]: probability_decision
```

```
In [62]: fpr_decision,tpr_decision,thresholds_decision = roc_curve(y_test,probability_decision)
```

```
In [63]: plt.plot(fpr_decision,tpr_decision)
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('ROC CURVE FOR DECISION TREE MODEL')
plt.show()
```

ROC CURVE FOR DECISION TREE MODEL

